

People & Organizational Performance Practice

# The future of work is agentic

The digital workforce is happening. Here's what it may look like when humans are working side by side with AI agents—and how to prepare now for this surprisingly near-term eventuality.



**Think about your** org chart. Now imagine it features both your current colleagues—humans, if you’re like most of us—and AI agents. That’s not science fiction; it’s happening—and it’s happening relatively quickly, according to McKinsey Senior Partner [Jorge Amar](#). In this episode of [McKinsey Talks Talent](#), Jorge joins McKinsey talent leaders [Brooke Weddle](#) and [Bryan Hancock](#) and Global Editorial Director Lucia Rahilly to talk about what these AI agents are, how they’re being used, and how leaders can prepare now for the workforce of the not-too-distant future.

The following transcript has been edited for clarity and length.

## From generative to agentic AI

**Lucia Rahilly:** Jorge, welcome to *McKinsey Talks Talent*.

**Jorge Amar:** Thank you very much. Excited to be here.

**Lucia Rahilly:** Jorge, there was a great little piece in *The Wall Street Journal* called “[Everyone's talking about AI agents. Barely anyone knows what they are.](#)” What exactly do we mean when we talk about agentic AI?

**Jorge Amar:** I'll start where I think most people still are, which is generative AI. Gen AI is mostly a reactive type of AI focused on generating creative content, triggered by a prompt or an instruction from an individual.

Now if we continue the evolution of AI into agentic, we start to come to a very different reality. The first difference is we're talking about AI that is not only generating content. It is executing on a task, on a mandate, on a particular instruction. An AI agent is perceiving reality based on its training. It then decides, applies judgment, and executes something. And that execution then reinforces its learning. It learns if what the agent did was good or bad and then feeds that back in.

So we're getting into the next step: AI deciding what to do on its own. We start to get into this complete AI workforce. Your AI agents could now be the evolution and the creation of a digital replica of the entire workforce of an organization.

**Lucia Rahilly:** OK, Jorge. You're scaring us. Let's talk through some use cases that might help bring this to life a bit. What does agentic AI look like now in the wild?

**Jorge Amar:** It's still the Wild, Wild West out there. But I'll try. Right now, many companies are starting to experiment. Typically, the environments in which they are deploying agents are very deterministic, with a clear process to follow. Think of IT help desks, or software development, or customer service tickets: any environment where a customer asks for something and there's a well-defined process afterward. The agent picks it up—decides what is the right process, the right content article to be retrieved, the right information to be gathered—and then triggers an action.

**‘Your AI agents could now be the evolution and the creation of a digital replica of the entire workforce of an organization.’**

—Jorge Amar

**Bryan Hancock:** In HR, we’re seeing agentic AI in talent acquisition. Agents clean records. They try to understand, “Of the vast universe of potential candidates, how do we clean the data and understand who the right candidate might be?” Then a separate agent goes through and scores those candidates and does the ranking and the sourcing process. A separate agent reaches out to gain contact and schedule interviews.

And then I’ve seen a coordinating agent that sits on top of the overall process, interacting with those underlying agents. Have you seen that kind of coordinating agent process? And how do you even create an agent that coordinates across some of those discrete subprocesses?

**Jorge Amar:** I have a client that is already doing the first screening of all candidates for the front line entirely with agents. And I have even seen one step further: AI agents being deployed for training.

Think of a call center or store environment. You generate an agentic customer, and you say, “This is a type of call. This is a type of customer. Record an interaction.” It’s not only simulating a real phone call. You’re also getting live, detailed scoring of how you, as a frontline employee, are doing in that interaction. Are you using the right words? Have you remembered every single step of the process? It gives you very detailed coaching instructions.

Probably before, a supervisor in a call center could listen to three, five calls per agent. Now you get a summary of every single call, with a detailed breakdown of all the things this human agent is doing well and could do better.

You can focus your coaching, your onboarding in a much more targeted way because you know exactly which skills to develop, which traits to emphasize. And it's not only recruiting and training; you could even do the same thing for [performance management](#).

## How agentic AI is already changing work

**Brooke Weddle:** Jorge, it sounds like you're pointing to examples where agentic AI has allowed companies to achieve greater levels of productivity. Recently, the [Work Trend Index annual report](#) came out, one of Microsoft's flagship publications in workforce. And it found that a third of executives are considering using AI to reduce head count in the next 12 to 18 months. But nearly 50 percent said they were considering maintaining head count but using AI as digital labor to boost productivity—as complementary to human skills. What have you seen in terms of use cases?

**Jorge Amar:** It is still early days when it comes to what I call decoupling the creation of capacity—automating tasks that otherwise would have been performed by a human—and the monetization of capacity. The monetization of capacity is its own independent thing.

One of the potential paths can be, "I'm going to reduce head count." More and more, some executives I'm talking to are interested in, "I might reduce head count, but I also might want to do things differently." Suppose your competitive advantage was your call center agents. If AI brings everyone to the same parity level, how do you differentiate? What are the implications for your workforce when you can differentiate by having the best algorithm, the best agentic framework out there—but at the same time, how do you complement that with humans to do things that otherwise would have been cost prohibitive?

I'll give you one example. Last week, I was talking with one of my travel clients—and you could pick the airline or cruise line of your choice. What if you now had your own personalized concierge looking at your travel, giving you very detailed recommendations on how to navigate the airport, suggesting the type of food you could pick up on your way and even creating the order for you, and then getting to the final point in which it helps you board the plane and makes sure you have space? The possibilities are endless when it comes to figuring out or creating new and different workflows, new processes, new ways to surprise and delight your customers that you couldn't have otherwise.

**Bryan Hancock:** And I imagine you can also do some of the same toward your employees. How do you surprise and delight across the employee journey? How do these agents actually get created—and get created in a way that's specific to processes in any one area?

**Jorge Amar:** We're all still figuring out the best way. There was a quote recently along the lines of, "IT will be the HR of AI agents of the future." I would divide the creation of an agent into a few different steps. This helps us understand who is doing what.

First, there clearly needs to be a rationale from the business: customer support, marketing, sales, HR. They would define, “What is the need for an AI capacity?” and decide, “What are the parameters of what this AI capability needs to perform?”

Then they would work with their IT or AI function to either develop or procure their agentic capabilities. In many cases, the specificity and complexity of these AI capabilities will require these companies to develop their agent capabilities in-house, because they cannot find them in the market. It’s going to be a hybrid situation.

Once that capability exists, you have to onboard and train that agent, which we call “tuning” an agent. Tuning the agent requires a number of things: a good articulation and understanding of the process you are trying to “agentize,” as well as a subject matter expert who really understands the ins and outs.

You also need someone who understands the available data—a content specialist who is saying, “These are the content articles, the corpus of knowledge you need to train your agent,” and who makes sure that knowledge is up to date. In one of my cases, we trained the agent, and the agent started to spit out a bunch of COVID-related policies that were no longer relevant. So you need to make sure the data is accurate, relevant, and up to date.

Last, you need a good, robust prompt-engineering skill set: someone who can teach, train, and tune the agent by saying, “When the customer or your employee says this, this is what they mean. This is what they are trying to accomplish. And therefore do X, Y, Z.”

## **Building and managing an AI workforce**

**Brooke Weddle:** Jorge, you mentioned [IT becoming the HR of AI agents](#). And, of course, it was Jensen Huang, the CEO of Nvidia, who said this recently. When you think about a digital workforce, whose job is it to ensure that digital workforce is reaching its full potential? Is it more in the realm of IT? Or is this a space where HR might have a few things to say, since for a long time, getting managers to reach their full potential has been more their purview?

**Jorge Amar:** Some pioneering companies in this space are expressing their org charts not only in number of FTEs [full-time employees] but also in number of agents being deployed in every part of the organization. So I think we are going into a world where you’ll have to think about your workforce as both agentic and human. And I don’t think IT will be able to do this alone. IT will be critical in enabling the foundational elements to train an agent—the data stack, the right procurement, the right platform for training and tuning the agents.

Now, the true missing pieces: one is the business. Nobody will be able to train an agent if you don’t know intimately the policies, the processes, what really differentiates you from a business perspective.

And then I think HR will play a key role—first, to really push the business on what can be done from a hybrid workforce perspective. Second, and we started seeing this in one of my clients, is where the technology is up and running, but the number of live interactions is not coming down. There is a big change management component that comes into play. HR will be absolutely critical there.

How do you tell your 20-year-tenured employee in the call center, “Now there is this agent that is going to do the job much better than you”? This person would probably say, “How can this AI thingy that got trained yesterday replace my 20 years of experience?” And there is a big step toward driving the incentives for usage, role modeling communication, and creating the right change story for these employees to understand, “Look at all the great possibilities this unlocks for you.”

This tells me HR still will play a critical role in the adoption of this agentic workforce. Maybe HR will not be screening each resume, but it will be critical in driving the change management efforts in adopting an agentic AI workforce.

**Lucia Rahilly:** Jorge, it’s so interesting to hear the anthropomorphic terms you’re using to describe these agents—existing in the org chart, for example, or as a digital workforce. To be clear, are these agents being construed as tools or as a class of digital workers—neo-colleagues of some kind?

**Jorge Amar:** I do think of it as a workforce. This is a workforce that will conduct end-to-end processes, replacing many tasks being performed today by the human workforce. It will augment the tasks a human workforce is performing to help make it better, faster, more efficient.

Some companies out there are even promoting this notion of a zero-FTE department—an entire function fully performed by an agent. Then you have on the side humans in the loop controlling or monitoring what these agents are doing. Putting philosophical debate aside, I think we should think of agents as a parallel workforce for all intents and purposes.

## Earning employee trust in the AI age

**Lucia Rahilly:** You mentioned adoption, and we so often hear adoption cited as a primary challenge in realizing the value of AI. How do you see humans in the workplace taking to this notion of collaborating with AI agents?

**Jorge Amar:** It’s still a big challenge. I’ll give you one example. In some of the frontline environments where I spend a lot of time, some of the newer agents or the newer reps tend to embrace AI faster. Why? Because if you’re just coming into a frontline environment, the back office is where you need to learn all these things, and now AI is guiding you through the process. That’s great. It makes the job easier. But some of the more tenured employees resist AI quite a bit. It’s really challenging for them.

# 'So I think we are going into a world where you'll have to think about your workforce as both agentic and human.'

—Jorge Amar

The other big element is that many employees tell us, "I cannot trust an AI black box out there that is doing this, so I will use the AI result, but at the same time, I'm going to have my own calculations." Therefore, you're now duplicating work. There are many of these elements that will be critical in cracking the code to adoption, because my fear is that we will end up with huge investments and very little value realized.

**Bryan Hancock:** Who do you think is going to lead the way in adoption?

**Jorge Amar:** First, there's got to be a clear mandate from the top. Leaders should make sure they are role modeling and integrating AI into the way they speak and what they do.

Second, evaluate the performance of AI in a joint fashion. One of my clients sees the results of both the human and the agentic parts of the operation in the same dashboard. The business manager, the VP, and the SVP evaluate the joint performance of both their workforces.

Third, this space is changing week by week, day by day. You need to design an operating model, a set of processes, that allows you to adapt. The more flexible this operating model, the better, because otherwise you're going to be making investments in a technology or a set of algorithms that three months from now are going to be different.

If you put all that in the mix, some of the smaller companies, start-up environments, have a little bit of an advantage. But the reality is that some of these LLMs [large language models] or agent platforms are not going to be trained on small companies. So it is critical to get to the larger companies and say, "Hey, I'm going to make the performance of these even better." How to do that in an effective way in that environment is, to me, the crux of this issue.

**Brooke Weddle:** What skills are going to become more salient in human leaders to get the most out of agents?

**'Some companies out there are even promoting this notion of a zero-FTE department—an entire function fully performed by an agent.'**

—Jorge Amar

**Jorge Amar:** First, HR will need to be at least business proficient in what an agentic workforce can do. How can you drive a change management program if you don't know what your agentic workforce can and cannot do, or what will be possible in three years?

Second, I think HR will play an important role in reskilling human employees. Today, you can probably fully agentize the workload of a level-one support engineer. But you might want to repurpose that person to become a prompt engineer or to do content generation for AI training. An HR function that can do that at scale is another critical component and skill set that HR will need to develop if you think about the next three, five years: "What is the evolution of that role?"

Last is being really good at [empathy](#), understanding the change story, helping employees onboard into their own AI journey, and making it happen in a way that is not threatening: "Look at all the other possibilities you might have in the future within the organization." Articulating that very clearly and helping employees come along in that journey is going to be another critical component.

**Brooke Weddle:** The *Work Trend Index annual report* I mentioned earlier talks about the need to evolve from an org chart to a work chart.

**Jorge Amar:** Yes, and you probably saw that the CEO of Shopify released a memo saying something along those lines: "Before you ask for new head count, show me that AI cannot do the work."

**Brooke Weddle:** That was almost positioned as a more radical stance. But in my conversations, it's very much a part of the conversation already. I very much think that's a now thing versus a future thing.

**Jorge Amar:** There are a couple of elements we also need to put on the table to say why now or not now. I would describe them in three broad categories. Number one is that to get an agent up

**'There are many of these elements that will be critical in cracking the code to adoption, because my fear is that we will end up with huge investments and very little value realized.'**

—Jorge Amar

and running, you do need a good technology stack and data stack. And there are many things being done to create new data, generate what we call synthetic data for training purposes.

Number two, there are a number of concerns about security and risks, from drift, hallucination, bias, and any of the challenges with some of these LLMs. For example, what if an agent is talking to your customer support agent, and they generate their own little dynamics and negotiation, and now suddenly you end up with a 90 percent discount on your product because you trained your agent into churn reduction and churn avoidance? How do you control that? Maybe you need to train a whole new set of agents that are monitoring the different negotiations and different discounts, and anything that touches your CRM [customer relationship management].

And the third is, "What is the cost? What are the different usability considerations from a UX [user experience] and UI [user interface] perspective?" It's great that you might have a very conversational chatbot, but if it looks like the 1990s interface of how you were interacting on some of your most famous messaging platforms, customers are not going to use it. So I think it is a very now conversation, but it also requires us to tackle some of these issues around risk, data, usability. Because otherwise, it's going to go into purgatory.

**Brooke Weddle:** That's not where we want to go. Clear.

**Lucia Rahilly:** Obviously, it's vital to be talking about this now, planning for it now. But acknowledging that predictions are freighted with uncertainty, what time frame do you think we're talking about for agents really to take effect at scale in companies?

**Jorge Amar:** It depends on who you ask. Some of the hyperscalers and technology companies would tell you that they are already deploying it, and they are. Many of the other organizations I talk to are saying, "I need to understand this; I need to evaluate it." And we're probably looking at 18 to 24 months out before it reaches full scale. I believe that there are a few elements where it'll take a little bit of time, making sure everyone is comfortable deploying them at scale.

## Preparing the workforce of tomorrow

**Bryan Hancock:** Jorge, I've got two college-age kids. What advice do you have for them as they're thinking through their careers and how to engage in work in a future that is agentic?

**Jorge Amar:** I was having this conversation a few weeks ago with a friend's son who asked me: "Maybe I should just drop out of college and become a prompt engineer." And look, I think there are certain jobs that are going to be fully transformed by AI. These net-new roles, such as prompt engineer and content specialist, will become more relevant in organizations. I would expect this demand to be higher than what the market can offer when it comes to just college. Therefore, I think we will have to go through a reskilling at scale within the existing workforce.

On the other side, how do you differentiate? If you differentiate only by having the best prompt engineer, fine. I think that is a skill set that at a certain point you will catch up on, because you could even have an agent that does prompt engineering.

But if you think the most important element a company has is the trust of and the relationship with their customers, do you need [a human workforce that is more empathetic](#)? Because, again, you might be OK talking to a chatbot to reschedule an appointment. But if you were just in an accident, do you want to talk to a human, or do you want to talk to a bot?

How do you emphasize skills in the incoming human workforce that help a company establish relationships? This could be the source of differentiation for your company. This could be the competitive advantage: "I offer a superior service. I offer a more human touch and surprise-and-delight experience."

For my friend's son, I was opening his mind in terms of, "Hey, maybe prompt engineering is fine, but maybe my arts background will be valuable in tomorrow's workforce because I will be able to understand human feelings in a way that no agent will be able to do."

**Brooke Weddle:** Jorge, if I reflect on what you're saying, I think it's a good time to consider the broader cultural implications of having a digital workforce. And some of that relates back to the values of a company. As you think about incorporating and onboarding agents as part of an organization, how do you do that in a way that is consistent with your company values, where you might prioritize collaboration, psychological safety, or having difficult conversations? It's a really interesting question to ask to get full value from the digital workforce.

# 'You might be OK talking to a chatbot to reschedule an appointment. But if you were just in an accident, do you want to talk to a human, or do you want to talk to a bot?'

—Jorge Amar

**Jorge Amar:** A hundred percent. That's why we are seeing more and more companies start to experiment with employee-facing agents more than just full end-customer-facing agents. Because how do you make sure that every interaction is in line with your corporate values, with your identity, with your brand standards, with the way you want to address a customer?

That's why I think we're going to go first through testing and scaling of an employee-facing agentic workforce. And then, over time, in certain discrete moments, you might want to do it with your end customer. You might want to do certain tasks that are mundane: customer authentication or verification or call summarization. But, again, you don't want to outsource to an agent or a digital agent the core of the relationship with your customer—or not just yet.

**Lucia Rahilly:** I read the article you recently coauthored, Jorge, about [agentic AI in the context of customer care](#). I found it fascinating that one of the findings was that almost three-quarters of Gen Z respondents to your survey believed live calls were quickest and simplest. Even younger cohorts seem to prefer human interaction. That points to the tremendous change management process that will have to happen for this to take hold.

**Jorge Amar:** So funny you mention that. Gen Zers would be bothered if they got a phone call from their parents, right? They would prefer to interact by text. I am sure anyone who has kids can relate to that. But for customer support needs, they prefer to talk.

So when we were digging a little deeper into why they prefer to talk to their provider, insurance company, telco carrier, bank, they all mentioned the same: "My situation is so unique, so important to me, that I just want to talk to a human who will give me that personalized and unique solution I won't be able to get through a bot." And the reality is maybe 80 percent, 90 percent of those interactions are the human giving them the exact same process, but still they felt they had a much more personalized solution by talking to a human.

Find more content like this on the  
**McKinsey Insights App**



Scan • Download • Personalize



So maybe that changes over time and customer support bots will just get true characters that enable a fully formed dialogue. To your point, the change management is not only with employees. It's also customer education and building trust on some of these solutions.

**Brooke Weddle:** Jorge, when you think about the next three to five years, what are you most optimistic and excited about when it comes to the potential of agentic?

**Jorge Amar:** I am really looking forward to doing things in a way that we couldn't have done otherwise. I am super optimistic about doing personalization at scale with customers. I am super optimistic about empowering humans to do tasks that are not repetitive, that are not going to create attrition levels of 50, 60, 100 percent per year, that create career paths for employees that are about connecting with humans—transforming the way we work on a daily basis, focusing on the change management elements we were just describing.

This hybrid workforce future should be a very uplifting environment for everyone—mostly for us as part of the workforce. It creates a new set of skills that were probably deprioritized in previous ways of finding efficiencies in companies—for example, trying to do things as repetitively and as fast as possible—and really opens the door to new ways of interacting with customers and employees.

**Bryan Hancock:** I have a fun question. Should we tell AI thank you?

**Jorge Amar:** I do because I think that when Skynet takes over, I want them to know I was very kind to them. It's funny. I was reading the other day that OpenAI is spending tons of cycles—I don't know if you've seen that news—on the use of "please" and "thank you."

**Brooke Weddle:** I say thank you to Alexa. It's just good behavior all around.

**Jorge Amar:** Of course.

**Bryan Hancock:** The article on OpenAI was saying, "We're spending millions of dollars, pumping tons of CO<sub>2</sub> into the atmosphere because of the energy used by the data centers because we're saying 'thank you.'"

**Jorge Amar:** Totally. But imagine your kids. You're going to teach them not to say thank you to Alexa but say thank you to a human? Come on.

**Jorge Amar** is a senior partner in McKinsey's Miami office and leads McKinsey's digital customer care group globally. **Brooke Weddle** is a senior partner in the Washington, DC, office, where **Bryan Hancock** is a partner. **Lucia Rahilly** is the global editorial director and deputy publisher of McKinsey Global Publishing and is based in the New York office.

Copyright © 2025 McKinsey & Company. All rights reserved.